Art Generation & Neural Style Transfer

Exploring Loss Functions and Evaluation Metrics in Deep Learning

Presented by: Mir Ahmed

Roll No: 23i-0142

Project Vision & Objective

Blending Content and Style

Neural Style Transfer (NST) is an optimization technique used to take three images: a **Content image**, a **Style Reference image**, and the **Target image** to produce a blended artwork.

- Maintains the structural integrity of the original photo.
- Applies artistic textures, strokes, and color palettes.
- Utilizes deep feature extraction from pre-trained models.

The Engine: Pre-trained VGG-19

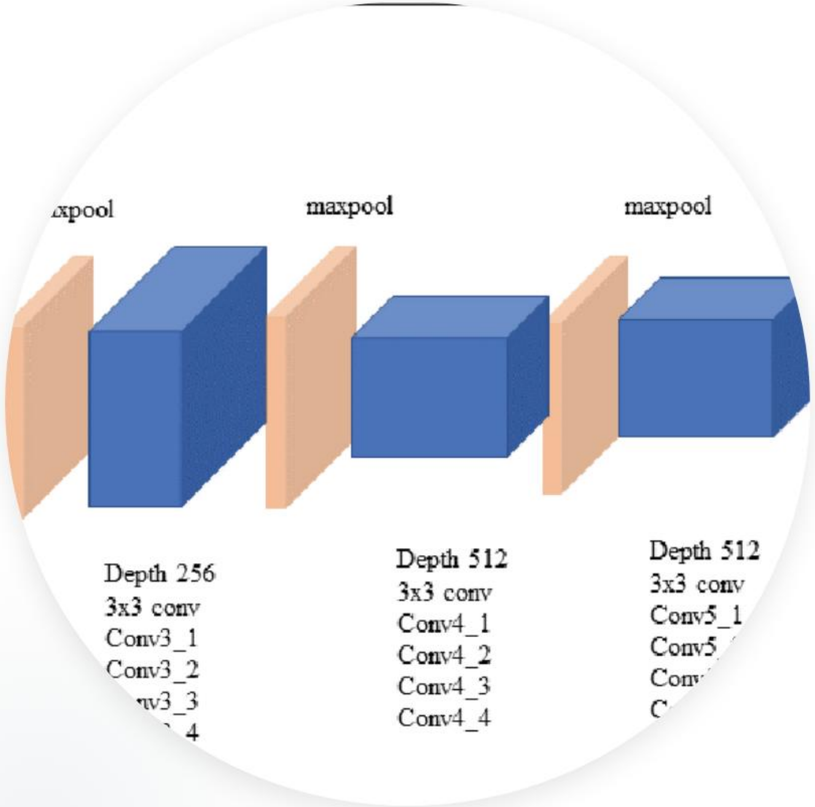

Convolutional Feature Extraction

We leverage the **VGG-19** network, trained on ImageNet, to understand image semantics at different depths:

- **Lower Layers:** Detect simple textures, edges, and colors (Style).
- **Higher Layers:** Capture complex spatial arrangements and global structure (Content).

Loss Functions

The mathematical heart of Neural Style Transfer optimization.

The Optimization Objective

NST works by minimizing a **Total Loss function**. This forces the generated image to simultaneously resemble the structure of the content and the feel of the style.

Total Loss Equation:

$$L_{total} = \alpha L_{content} + \beta L_{style}$$

Content Loss ($L_{content}$)

Calculates Euclidean distance between the feature maps of the content and generated images.

Style Loss (L_{style})

Measures the difference in texture correlation using **Gram Matrices** across multiple layers.

Perceptual Loss vs. MSE

Perceptual (VGG) Loss

Focuses on **High-Level Semantics**. It compares images based on their activations in a pre-trained network.

Result: Sharp, meaningful artistic textures that "understand" the context.

Standard MSE Loss

Focuses on **Pixel-Level Values**. It measures the raw reconstruction error at the individual pixel coordinate level.

Result: Often blurry or ghost-like outputs as it struggles with high-frequency artistic details.

Comparison of Loss Functions

Metric / Loss	Advantages	Limitations	Best Use Case
Perceptual Loss	Captures high-level textures & semantic patterns.	Computationally heavy; needs pre-trained CNN.	Style Transfer & Super Resolution.
MSE Loss	Extremely simple; fast to calculate.	Inaccurate for human perception; blurry art.	Simple image reconstruction tasks.
Style Loss	Excellent at capturing artistic brush strokes.	Sensitive to weight hyper-parameters (α, β).	Art Generation & Neural Painting.

Evaluation Metrics

Quantifying visual quality, realism, and structural fidelity.

Fréchet Inception Distance (FID)

FID

REALISM & DIVERSITY

Measuring Image Distribution

FID is the gold standard for evaluating generated art. It compares the distribution of generated images with the distribution of real "style" paintings.

- **Low Score:** High similarity to real dataset distributions.
- **High Score:** Indicates "fake" looking or low-diversity art.
- **Robustness:** More reliable than Inception Score for art tasks.

SSIM and PSNR

SSIM

Structural Similarity Index
Measures degradation of luminance, contrast, and structure. Crucial for ensuring the user can still recognize the **Content** after stylization.

PSNR

Peak Signal-to-Noise Ratio
A classic engineering metric for image quality. It measures the ratio between the maximum power of a signal and the power of corrupting noise.

Human Perception

While SSIM is closer, neither perfectly captures **Artistic Value**. They must be used alongside qualitative expert review.

Summary of Evaluation Metrics

Metric	What it Measures	Advantages	Limitations
FID	Realism & Distribution distance.	Best for GANs/Art generation.	Needs large sample size to be stable.
SSIM	Structural Preservation.	Correlates well with human eye.	Fails to measure "Artistic" style fit.
PSNR	Pixel-wise Image Quality.	Easy to interpret and compute.	Ignores high-level artistic features.

Thank You!

Questions or Discussion?

Mir Ahmed | 23i-0142

Project: Neural Style Transfer & Art Generation

Image Sources

https://framerusercontent.com/images/rhhsWQyrUWrNuRdtTOqbvMFqL8.png?width=900&height=507

Source: www.v7labs.com

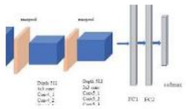

https://miro.medium.com/1*Q_bg1E3trWcjdk9_jnVGwg.png

Source: medium.com